evidence quality

25 questions/cell; quality/evidence: repetition 1; latency: all 75 measurements.

Intervals: descriptive, unadjusted 95% source-cluster bootstrap; no winner selection.

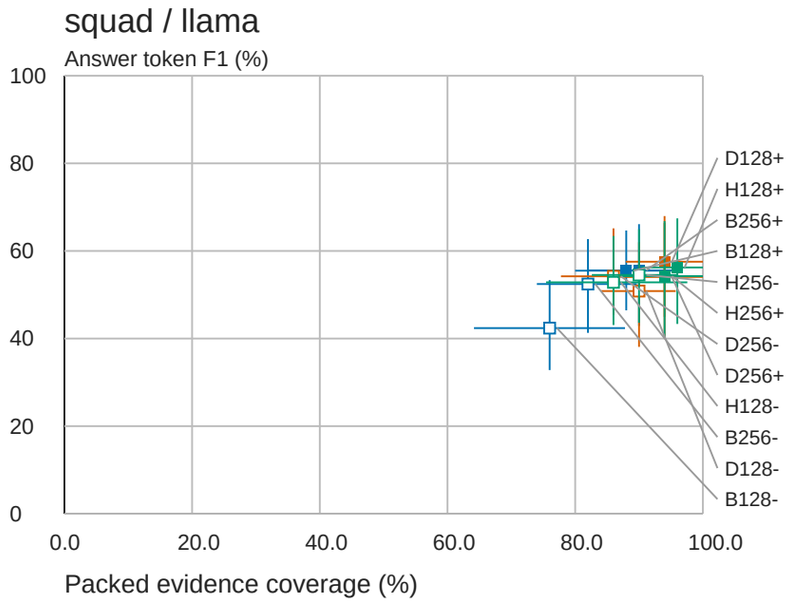

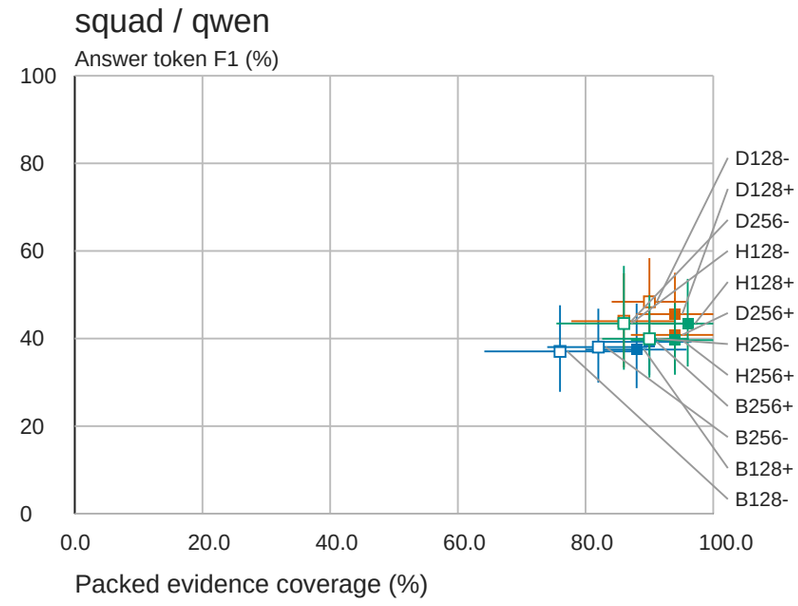

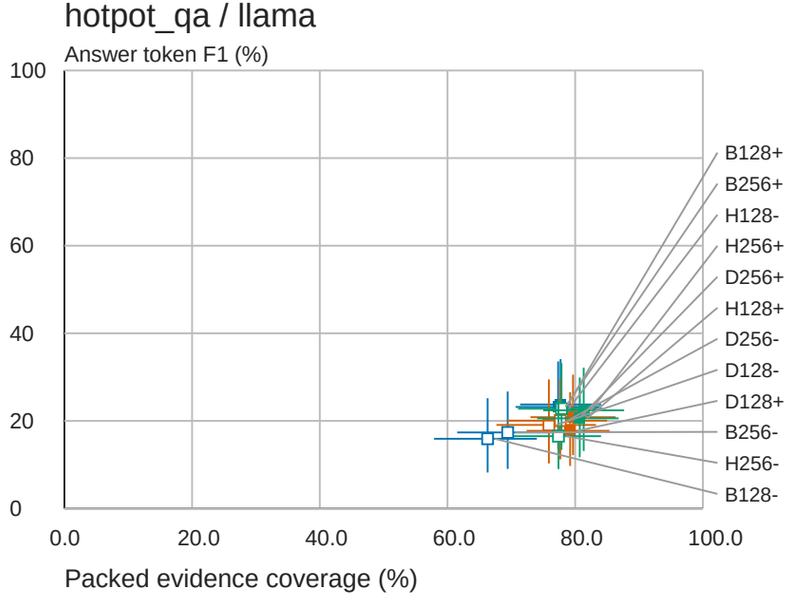

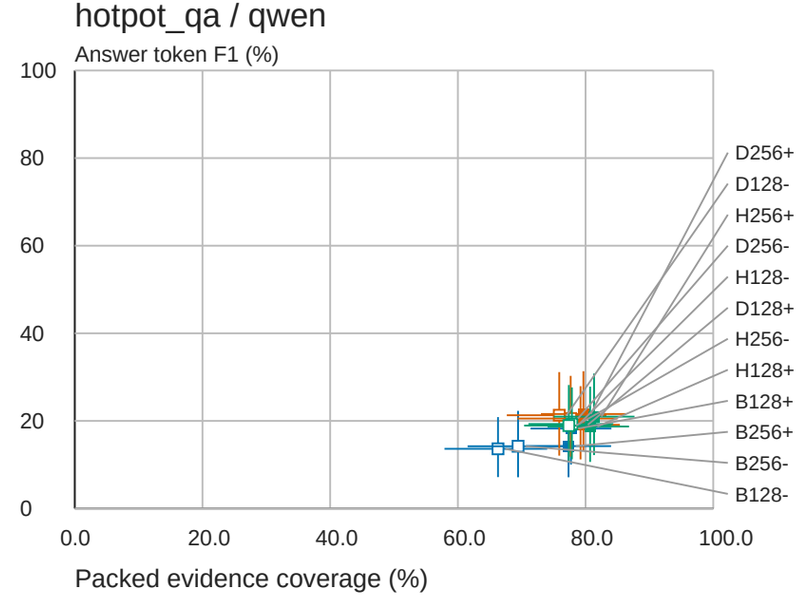

B=BM25, D=Dense, H=Hybrid; 128/256=chunk size; -=reranking off, +=on (filled).

Baseline plots separate: no-RAG keeps grounding instructions; gold is diagnostic evidence.

Baseline/main online implementations differ; do not pool their timings.